



ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA
CAMPUS DI FORLÌ

91258 / B0385
Natural Language Processing
Lesson 15. Convolutions in Text

Alberto Barrón-Cedeño
a.barron@unibo.it

27/11/2024

Table of Contents

1. Quick Keras Reminder
2. Prologue to CNN and RNN
3. CNN
4. CNNs for NLP

Chapter 7 of Lane et al. (2019)

Quick Keras Reminder

Keras

Sequential()

- Python class
- Neural network abstraction
- Grants access to the basic Keras API

Keras

Sequential()

- Python class
- Neural network abstraction
- Grants access to the basic Keras API

Sequential.compile()

- Builds the underlying weights
- Builds the interconnected relationships

Keras

Sequential()

- Python class
- Neural network abstraction
- Grants access to the basic Keras API

Sequential.compile()

- Builds the underlying weights
- Builds the interconnected relationships

Sequential.fit()

- Computes the training errors (loss)
- Applies backpropagation (weight adjustment)

Keras

Sequential()

- Python class
- Neural network abstraction
- Grants access to the basic Keras API

Sequential.compile()

- Builds the underlying weights
- Builds the interconnected relationships

Sequential.fit()

- Computes the training errors (loss)
- Applies backpropagation (weight adjustment)

Some “cooking” hyperparameters

`epochs` number of iterations over the data

`batch_size` number of instances before adjusting

`optimizer` function

Prologue to CNN and RNN

Prologue

- We have learned to build embedding spaces for words and texts

Prologue

- We have learned to build embedding spaces for words and texts
- We are considering the neighborhood of the words (~the bag)

Prologue

- We have learned to build embedding spaces for words and texts
- We are considering the neighborhood of the words (\sim the bag)
- We are not considering *actual* connections yet

Prologue

- We have learned to build embedding spaces for words and texts
- We are considering the neighborhood of the words (\sim the bag)
- We are not considering *actual* connections yet
- The downstream application is usually classification or regression

Prologue

- We have learned to build embedding spaces for words and texts
- We are considering the neighborhood of the words (\sim the bag)
- We are not considering *actual* connections yet
- The downstream application is usually classification or regression

We will start heading towards text generation

Words have relations and influence each other

Word order

$s_1 = \text{The dog chased the cat.}$

$s_2 = \text{The cat chased the dog.}$

(Lane et al., 2019, p. 220)

Words have relations and influence each other

Word order

$s_1 = \text{The dog chased the cat.}$

$s_2 = \text{The cat chased the dog.}$

$$\text{sim}(\text{tfidf}(s_1), \text{tfidf}(s_2)) = 1$$

Words have relations and influence each other

Word order

$s_1 = \text{The dog chased the cat.}$

$s_2 = \text{The cat chased the dog.}$

$$\text{sim}(\text{tfidf}(s_1), \text{tfidf}(s_2)) = 1$$

$$\text{sim}(\text{wv}(s_1), \text{wv}(s_2)) = 1$$

(1)

Words have relations and influence each other

Word order

s_1 = The dog chased the cat.

s_2 = The cat chased the dog.

$$\text{sim}(\text{tfidf}(s_1), \text{tfidf}(s_2)) = 1$$

$$\text{sim}(\text{wv}(s_1), \text{wv}(s_2)) = 1$$

(1)

But s_1 and s_2 are not the same!

Words have relations and influence each other

Word order

s_1 = The dog chased the cat.

s_2 = The cat chased the dog.

$$\text{sim}(\text{tfidf}(s_1), \text{tfidf}(s_2)) = 1$$

$$\text{sim}(\text{wv}(s_1), \text{wv}(s_2)) = 1$$

(1)

But s_1 and s_2 are not the same!

Word proximity

s = His mother, besides her son's willingness to amend the issue,
decided to punish him

Words have relations and influence each other

Word order

$s_1 = \text{The dog chased the cat.}$

$s_2 = \text{The cat chased the dog.}$

$$\text{sim}(\text{tfidf}(s_1), \text{tfidf}(s_2)) = 1$$

$$\text{sim}(\text{wv}(s_1), \text{wv}(s_2)) = 1$$

(1)

But s_1 and s_2 are not the same!

Word proximity

$s = \text{His mother, besides her son's willingness to amend the issue, decided to punish him}$

mother... decided | son... him

(Lane et al., 2019, p. 220)

Words have relations and influence each other

Spatial relation

Consider the position of words
(~written)

(Lane et al., 2019, p. 220)

Words have relations and influence each other

Spatial relation

Consider the position of words
(~written)

Temporal relation

Consider words as time series
(~spoken)

(Lane et al., 2019, p. 220)

Words have relations and influence each other

Spatial relation

Consider the position of words
(~written)

→ fixed-width window

Temporal relation

Consider words as time series
(~spoken)

Words have relations and influence each other

Spatial relation

Consider the position of words
(~written)

Temporal relation

Consider words as time series
(~spoken)

→ fixed-width window
convolutional neural networks

(Lane et al., 2019, p. 220)

Words have relations and influence each other

Spatial relation

Consider the position of words
(~written)

→ fixed-width window
convolutional neural networks

Temporal relation

Consider words as time series
(~spoken)

→ ongoing (unk) amount of time

Words have relations and influence each other

Spatial relation

Consider the position of words
(~written)

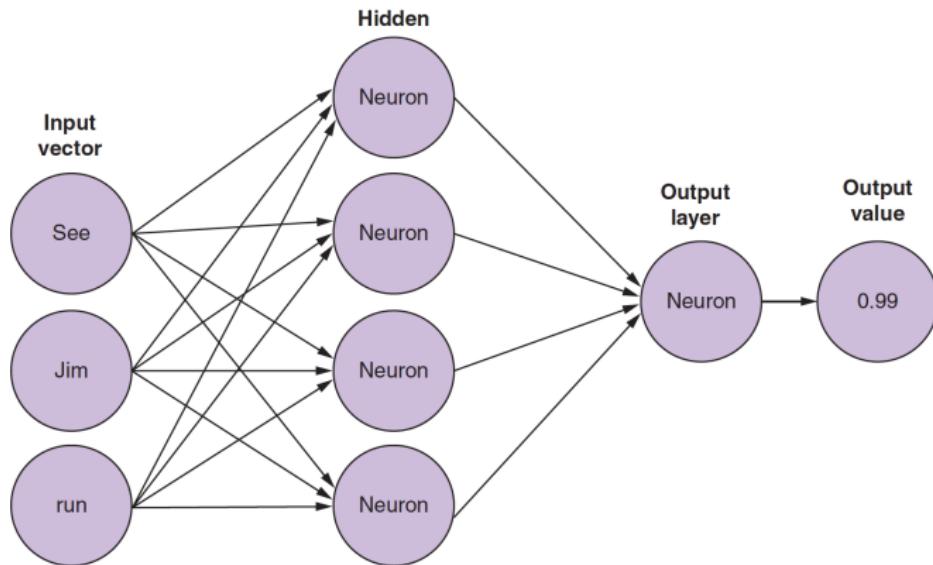
→ fixed-width window
convolutional neural networks

Temporal relation

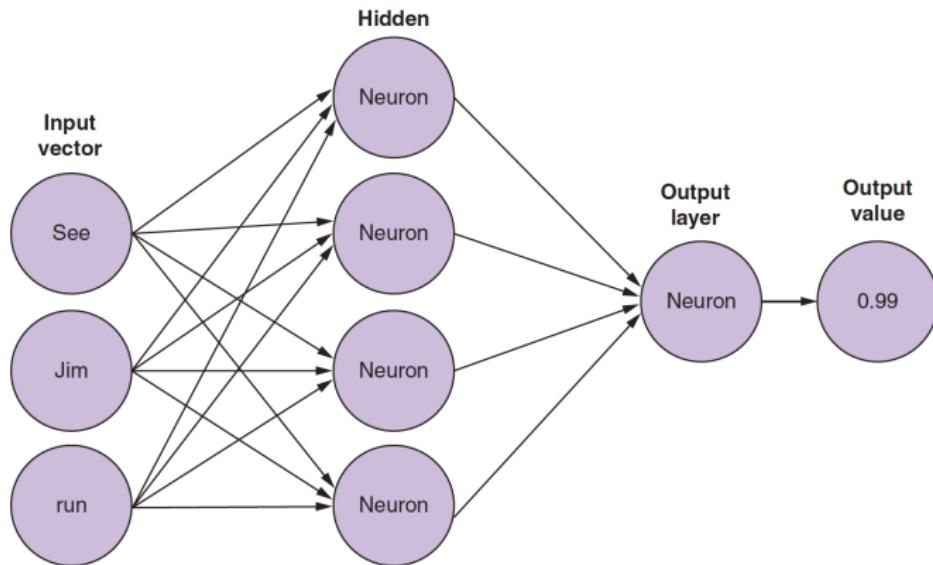
Consider words as time series
(~spoken)

→ ongoing (unk) amount of time
recurrent neural networks

Multiple Input Words

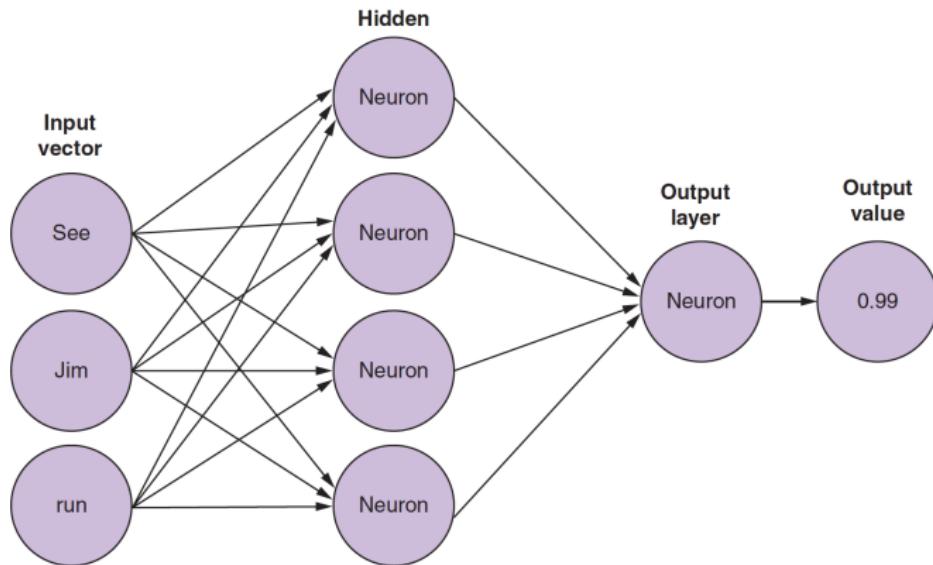


Multiple Input Words



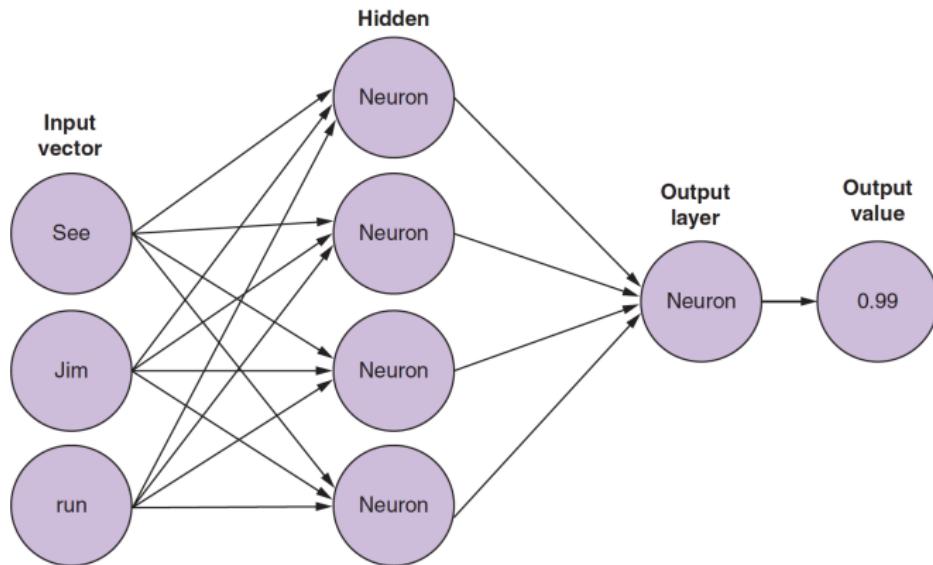
- Three tokens are passed at a time

Multiple Input Words



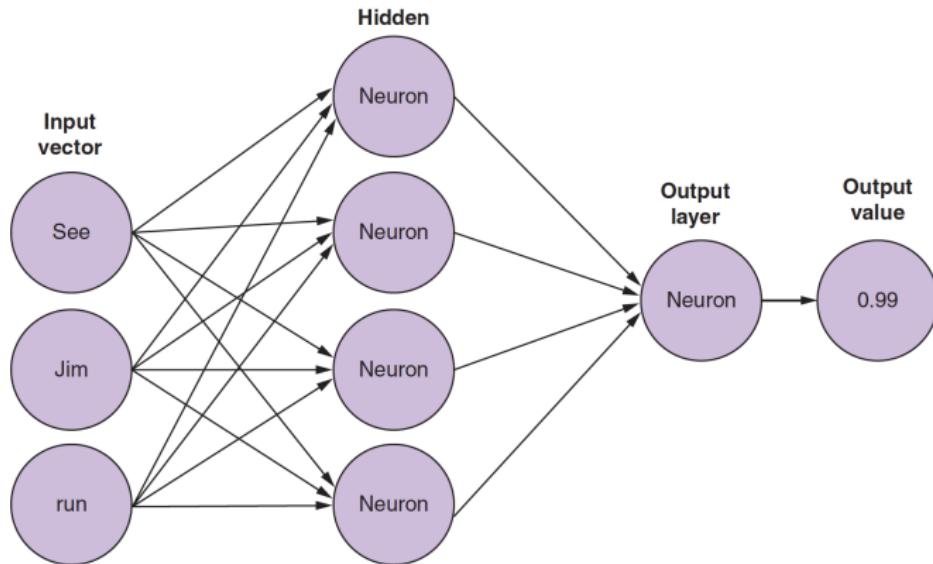
- Three tokens are passed at a time
- Two input alternatives
 - one-hot vector

Multiple Input Words



- Three tokens are passed at a time
- Two input alternatives
 - one-hot vector
 - pre-trained word vector

Multiple Input Words



- Three tokens are passed at a time
- Two input alternatives
 - one-hot vector
 - pre-trained word vector

See Jim run \neq run See Jim (!)

CNN

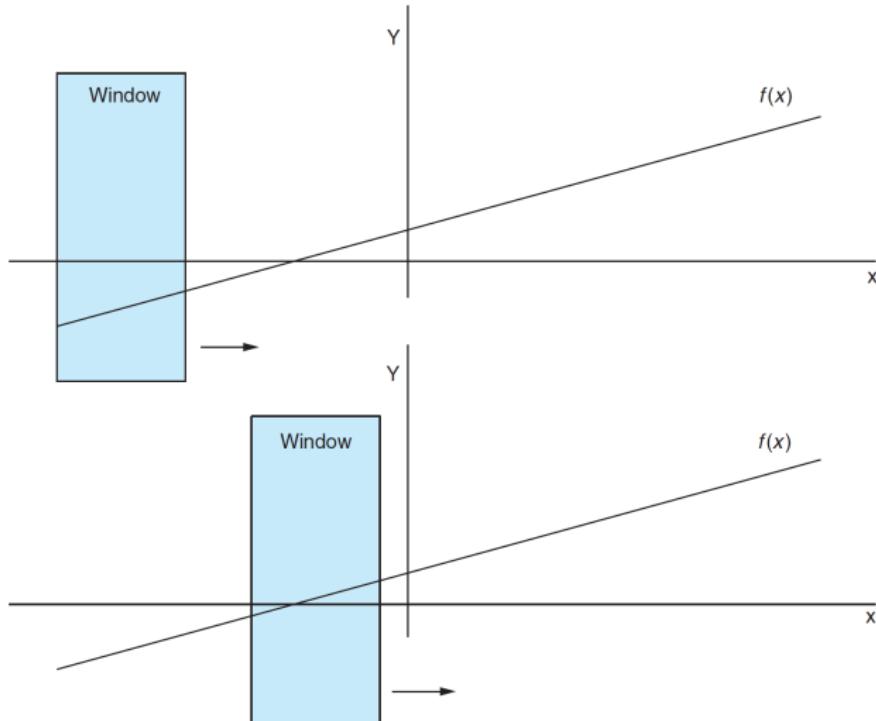
Convolutional Neural Networks

Sliding —or convolving¹— a window over the sample

¹To roll or wind together (Webster's)
(Lane et al., 2019, p. 222)

Convolutional Neural Networks

Sliding —or convolving¹— a window over the sample



¹To roll or wind together (Webster's)
(Lane et al., 2019, p. 222)

Convolutional Neural Networks

Back to the roots: image recognition

- Input: pixels of an image
- Output: the image contains x

[https:](https://blogs.nvidia.com/wp-content/uploads/2019/04/ADAS-IMG_0052.jpg)

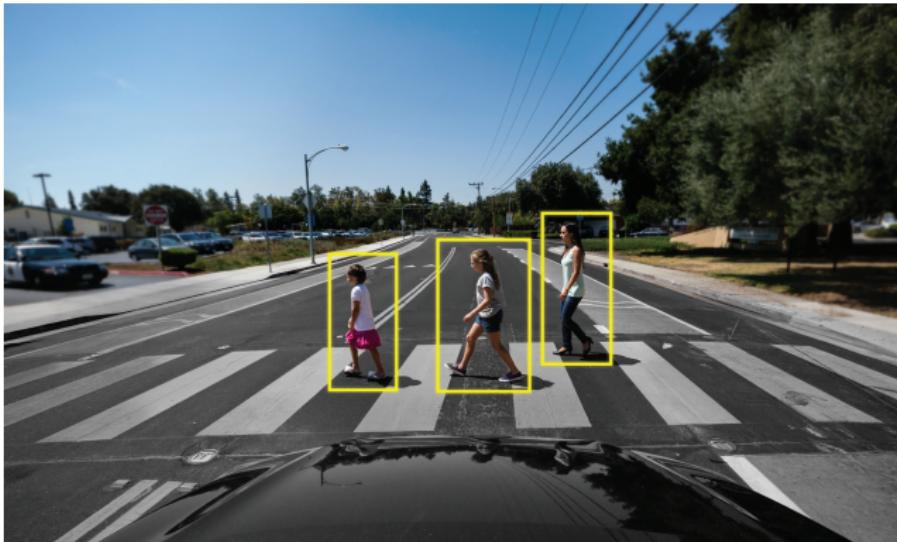
//blogs.nvidia.com/wp-content/uploads/2019/04/ADAS-IMG_0052.jpg



Convolutional Neural Networks

Back to the roots: image recognition

- Input: pixels of an image
- Output: the image contains x



https://blogs.nvidia.com/wp-content/uploads/2019/04/ADAS-IMG_0052.jpg

Convolutional Neural Networks

When the input is an image

- B&W: [0,1] (with a smooth binariser)
- Grayscaled: [0, 255]
- Colour: R: [0, 255] G: [0, 255] B: [0, 255]

Convolutional Neural Networks

When the input is an image

- B&W: [0,1] (with a smooth binariser)
- Grayscaled: [0, 255]
- Colour: R: [0, 255] G: [0, 255] B: [0, 255]



(Lane et al., 2019, p. 223)

Convolutional Neural Networks

When the input is an image

- B&W: [0,1] (with a smooth binariser)
- Grayscaled: [0, 255]
- Colour: R: [0, 255] G: [0, 255] B: [0, 255]



	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27			
0	120	119	118	108	103	111	113	111	117	120	120	119	119	121	114	118	120	120	120	121	121	121	115	100	120	118	117				
1	111	109	111	106	108	118	120	110	106	117	120	119	119	121	114	118	119	118	119	118	114	119	109	102	122	114	117				
2	109	115	121	116	108	118	118	104	111	115	119	119	119	119	118	110	121	122	123	73	92	128	109	87	119	115	113				
3	105	102	114	104	100	108	109	109	108	109	106	106	106	106	106	106	106	106	106	106	106	106	106	106	106	106	106				
4	108	110	100	116	111	95	104	110	114	117	117	117	117	117	117	106	107	109	94	90	42	70	62	86	70	96	118				
5	103	112	109	98	100	93	111	112	115	113	117	117	117	117	117	116	116	112	118	116	114	110	41	81	110	102	119	114	119	117	
6	111	111	112	104	93	106	119	115	108	110	116	115	116	117	115	107	108	111	118	49	87	105	100	115	110	115	115	115			
7	107	111	111	104	104	104	104	104	104	104	104	104	104	104	104	104	104	104	104	104	104	104	104	104	104	104	104	104			
8	112	105	108	94	99	102	95	106	111	113	111	111	108	106	109	112	109	114	112	114	115	45	83	104	108	111	112	114	114		
9	112	105	108	86	95	103	104	103	106	104	108	109	114	114	111	108	112	112	113	114	43	80	104	111	109	113	114	113			
10	99	111	102	88	88	111	107	101	101	106	111	112	112	110	113	111	107	112	112	113	112	113	112	113	110	108	114	112	113		
11	104	104	104	98	98	107	107	107	107	107	107	107	107	107	107	107	107	107	107	107	107	107	107	107	107	107	107	107	107		
12	101	93	76	101	103	107	107	108	110	107	103	111	111	111	110	109	106	112	111	111	111	109	27	82	108	107	111	112	111	100	
13	98	92	99	115	108	108	111	106	100	98	106	108	109	110	107	106	109	108	107	107	37	78	106	103	106	106	108	103	98		
14	100	73	97	102	92	95	93	89	89	97	103	103	106	106	102	101	106	104	106	103	33	75	105	103	108	109	98	107	107		
15	94	99	92	95	95	95	95	95	95	95	95	95	95	95	95	95	95	95	95	95	95	95	95	95	95	95	95	95	95	95	
16	71	82	87	85	78	89	106	104	99	106	106	105	106	105	105	105	105	105	105	105	105	107	103	21	72	106	106	109	100	102	109
17	67	87	64	68	84	89	98	96	99	104	104	104	105	103	101	105	103	104	102	23	76	103	103	106	98	108	101	101	101		
18	68	87	87	92	88	84	84	90	98	98	102	103	103	103	103	103	103	103	103	103	103	103	103	103	103	103	103	103	103		
19	60	91	87	77	80	85	82	82	82	82	82	82	82	82	82	82	82	82	82	82	82	82	82	82	82	82	82	82	82	79	
20	84	98	87	94	101	101	101	101	101	101	100	101	103	98	98	100	96	97	87	87	12	71	100	97	93	105	91	101	101		
21	92	80	88	99	100	98	100	100	97	98	97	97	96	92	91	92	93	96	87	13	79	105	95	98	96	89	100	100	100	100	
22	77	80	88	92	96	97	95	95	93	92	92	92	91	94	96	95	98	89	93	12	79	103	91	100	89	98	104	104	104	104	
23	71	84	84	84	84	84	84	84	84	84	84	84	84	84	84	84	84	84	84	84	84	74	88	88	88	88	88	88	88	88	88
24	60	66	82	92	90	90	87	90	94	94	94	93	94	92	91	93	94	95	95	81	0	76	90	92	81	77	65	58	58	58	
25	87	81	83	86	87	84	90	92	92	92	92	92	92	91	92	91	92	77	4	73	91	92	81	95	96	95	95	95	95		
26	87	88	88	83	85	89	91	89	90	91	91	91	91	90	92	90	89	73	8	66	92	83	84	92	91	91	91	91	91		
27	81	86	88	91	89	89	88	89	89	89	89	89	89	84	83	79	80	87	74	0	60	89	77	90	91	90	91	90	91	91	

(Lane et al., 2019, p. 223)

Convolutional Neural Networks

When the input is an image

An image is just a bunch of numbers

Convolutional Neural Networks

When the input is an image

An image is just a bunch of numbers

- Appropriate as input for a NN

Convolutional Neural Networks

When the input is an image

An image is just a bunch of numbers

- Appropriate as input for a NN
- But one single pixel has no real meaning

Convolutional Neural Networks

When the input is an image

An image is just a bunch of numbers

- Appropriate as input for a NN
- But one single pixel has no real meaning

→ Sliding over fragments of the image

Convolutional Neural Networks

When the input is an image

An image is just a bunch of numbers

- Appropriate as input for a NN
- But one single pixel has no real meaning

→ Sliding over fragments of the image

The convolution defines a set of filters (aka kernels) to do just that

Convolutional Neural Networks

When the input is an image

An image is just a bunch of numbers

- Appropriate as input for a NN
- But one single pixel has no real meaning

→ Sliding over fragments of the image

The convolution defines a set of filters (aka kernels) to do just that

- Take “snapshots” of different areas of the image
- Process them, one at a time

Convolutional Neural Networks

Strides and filters

Stride

- The distance “traveled” when sliding
- Yet another parameter
- Never bigger than the size of the filter → overlapping areas

Convolutional Neural Networks

Strides and filters

Stride

- The distance “traveled” when sliding
- Yet another parameter
- Never bigger than the size of the filter → overlapping areas

Sounds familiar?

Convolutional Neural Networks

Strides and filters

Stride

- The distance “traveled” when sliding
- Yet another parameter
- Never bigger than the size of the filter → overlapping areas

Sounds familiar? *n*-grams!

Convolutional Neural Networks

Strides and filters

Stride

- The distance “traveled” when sliding
- Yet another parameter
- Never bigger than the size of the filter → overlapping areas

Sounds familiar? *n*-grams!

Filter

- $n \times m$ surfaces

Convolutional Neural Networks

Strides and filters

Stride

- The distance “traveled” when sliding
- Yet another parameter
- Never bigger than the size of the filter → overlapping areas

Sounds familiar? *n*-grams!

Filter

- $n \times m$ surfaces
- Typically $n = m = 3$ (often $n \neq m$)

Convolutional Neural Networks

Strides and filters

Stride

- The distance “traveled” when sliding
- Yet another parameter
- Never bigger than the size of the filter → overlapping areas

Sounds familiar? *n*-grams!

Filter

- $n \times m$ surfaces
- Typically $n = m = 3$ (often $n \neq m$)
- Includes a set of weights (fix for the whole image)

Convolutional Neural Networks

Strides and filters

Stride

- The distance “traveled” when sliding
- Yet another parameter
- Never bigger than the size of the filter → overlapping areas

Sounds familiar? *n*-grams!

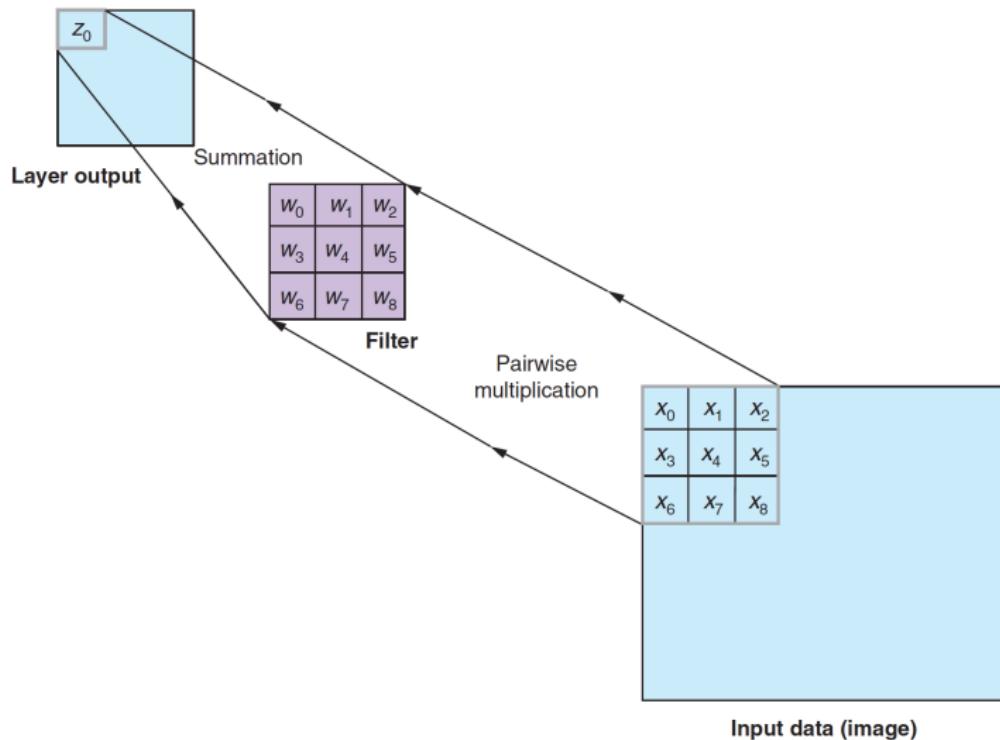
Filter

- $n \times m$ surfaces
- Typically $n = m = 3$ (often $n \neq m$)
- Includes a set of weights (fix for the whole image)
- Includes an activation function: usually ReLU

$$z = \max(\text{sum}(x * w), 0)$$

Convolutional Neural Networks

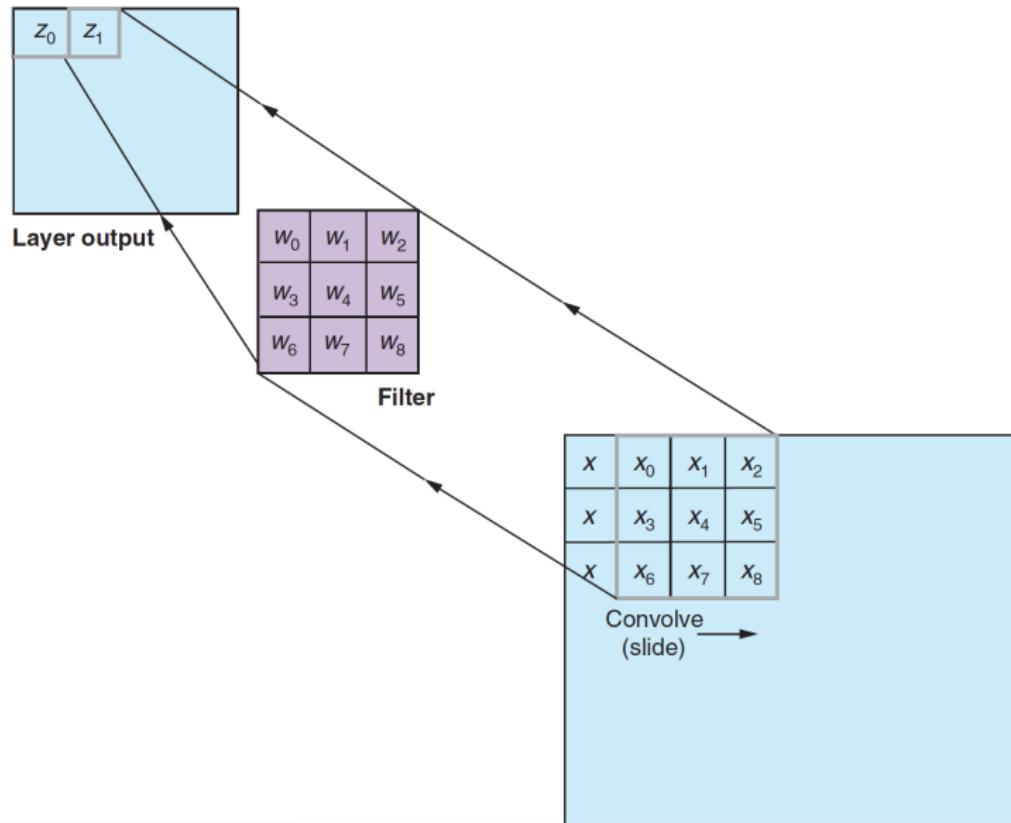
Convolutional step



(Lane et al., 2019, p. 225)

Convolutional Neural Networks

Convolution



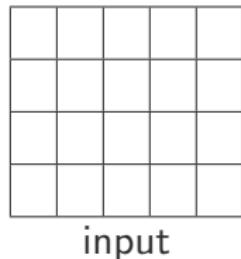
Convolutional Neural Networks

Producing multiple images

- k filters exist which carry out different operations
- Every filter will produce a new image, combination of source and filter

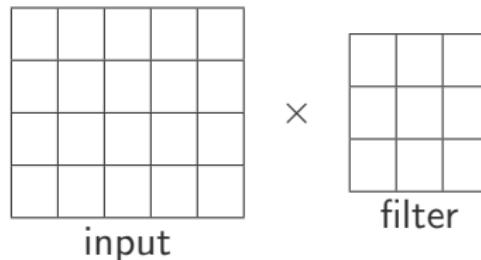
Convolutional Neural Networks

Padding



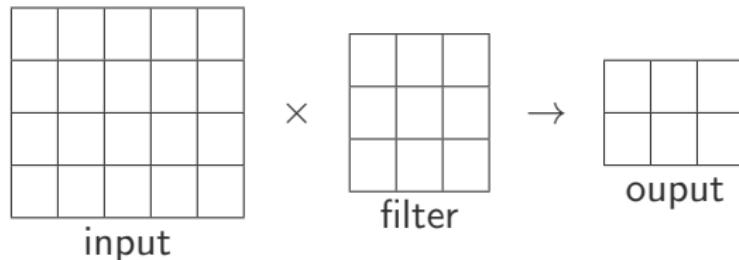
Convolutional Neural Networks

Padding



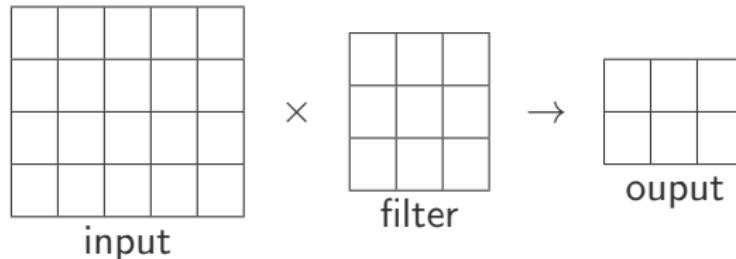
Convolutional Neural Networks

Padding



Convolutional Neural Networks

Padding

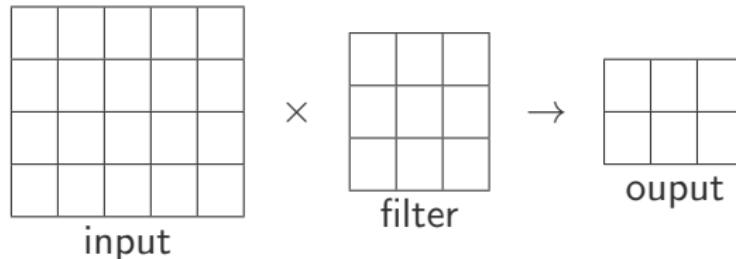


We are producing smaller images

"I don't care": Keras' argument `padding='valid'`

Convolutional Neural Networks

Padding



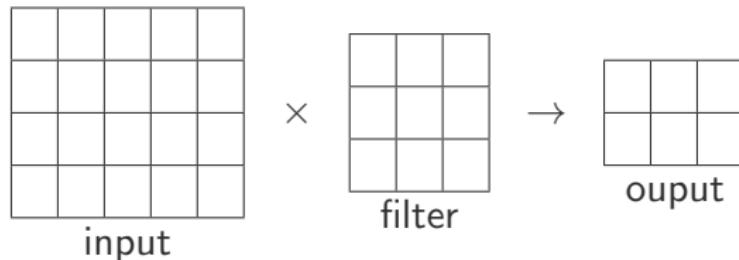
We are producing smaller images

"I don't care": Keras' argument `padding='valid'`

The edges of the image are undersampled

Convolutional Neural Networks

Padding



We are producing smaller images

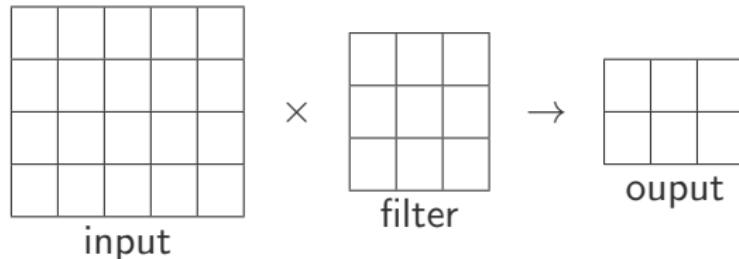
“I don't care”: Keras' argument padding='valid'

The edges of the image are undersampled

“I do care”: Keras' padding argument padding='same'

Convolutional Neural Networks

Padding



We are producing smaller images

"I don't care": Keras' argument padding='valid'

The edges of the image are undersampled

"I do care": Keras' padding argument padding='same'

In NLP we care

Convolutional Neural Networks

Pipeline

Input: an image, text

Output: a class, a real number

- Produce k new images through k filters

Convolutional Neural Networks

Pipeline

Input: an image, text

Output: a class, a real number

- Produce k new images through k filters
- Wire the filtered images to a feed-forward network

Convolutional Neural Networks

Pipeline

Input: an image, text

Output: a class, a real number

- Produce k new images through k filters
- Wire the filtered images to a feed-forward network
- Proceed as usual

Convolutional Neural Networks

Pipeline

Input: an image, text

Output: a class, a real number

- Produce k new images through k filters
- Wire the filtered images to a feed-forward network
- Proceed as usual

We can add multiple convolution layers

A full path of learning layers and abstractions

Convolutional Neural Networks

Pipeline

Input: an image, text

Output: a class, a real number

- Produce k new images through k filters
- Wire the filtered images to a feed-forward network
- Proceed as usual

We can add multiple convolution layers

A full path of learning layers and abstractions

- Edges

Convolutional Neural Networks

Pipeline

Input: an image, text

Output: a class, a real number

- Produce k new images through k filters
- Wire the filtered images to a feed-forward network
- Proceed as usual

We can add multiple convolution layers

A full path of learning layers and abstractions

- Edges
- Shapes

Convolutional Neural Networks

Pipeline

Input: an image, text

Output: a class, a real number

- Produce k new images through k filters
- Wire the filtered images to a feed-forward network
- Proceed as usual

We can add multiple convolution layers

A full path of learning layers and abstractions

- Edges
- Shapes
- Colours

Convolutional Neural Networks

Pipeline

Input: an image, text

Output: a class, a real number

- Produce k new images through k filters
- Wire the filtered images to a feed-forward network
- Proceed as usual

We can add multiple convolution layers

A full path of learning layers and abstractions

- Edges
- Shapes
- Colours
- Concepts

Convolutional Neural Networks

Pipeline

Input: an image, text

Output: a class, a real number

- Produce k new images through k filters
- Wire the filtered images to a feed-forward network
- Proceed as usual

We can add multiple convolution layers

A full path of learning layers and abstractions

- Edges
- Shapes
- Colours
- Concepts

What is learned

- Good filters
- “Standard” weights

Convolutional Neural Networks

Keras premier

```
from keras.models import Sequential  
from keras.layers import Conv1D  
  
model = Sequential()  
  
model.add(Conv1D(filters=16,  
                 kernel_size=3,  
                 padding='same',  
                 activation='relu',  
                 strides=1,  
                 input_shape=(100, 300))  
)
```

CNN Wrap up

- Sliding —or convolving— a window over the sample

CNN Wrap up

- Sliding —or convolving— a window over the sample
- Filters (kernels; matrices) slide over fragments of the image

CNN Wrap up

- Sliding —or convolving— a window over the sample
- Filters (kernels; matrices) slide over fragments of the image
- “Snapshots” of different areas of the image are taken and processed

CNN Wrap up

- Sliding —or convolving— a window over the sample
- Filters (kernels; matrices) slide over fragments of the image
- “Snapshots” of different areas of the image are taken and processed
- Multiple filters produce multiple images

CNN Wrap up

- Sliding —or convolving— a window over the sample
- Filters (kernels; matrices) slide over fragments of the image
- “Snapshots” of different areas of the image are taken and processed
- Multiple filters produce multiple images
- Multiple convolution layers can be added

CNN Wrap up

- Sliding —or convolving— a window over the sample
- Filters (kernels; matrices) slide over fragments of the image
- “Snapshots” of different areas of the image are taken and processed
- Multiple filters produce multiple images
- Multiple convolution layers can be added
- At the end, we can plug a “standard” fully-connected NN

CNNs for NLP

Back to Text

- In images both vertical and horizontal relationships are relevant

²|l2r or r2l; for some languages it's the vertical direction that matters (e.g., Japanese) (Lane et al., 2019, p. 229)

Back to Text

- In images both vertical and horizontal relationships are relevant
- In text only horizontal ones do²

²|l2r or r2l; for some languages it's the vertical direction that matters (e.g., Japanese) (Lane et al., 2019, p. 229)

Back to Text

- In images both vertical and horizontal relationships are relevant
- In text only horizontal ones do²
- We need “1D” filters

1 × 3 Filter

The cat and dog went to the bodega together.

1 × 3 Filter

The cat and dog went to the bodega together.

1 × 3 Filter

The cat and dog went to the bodega together.

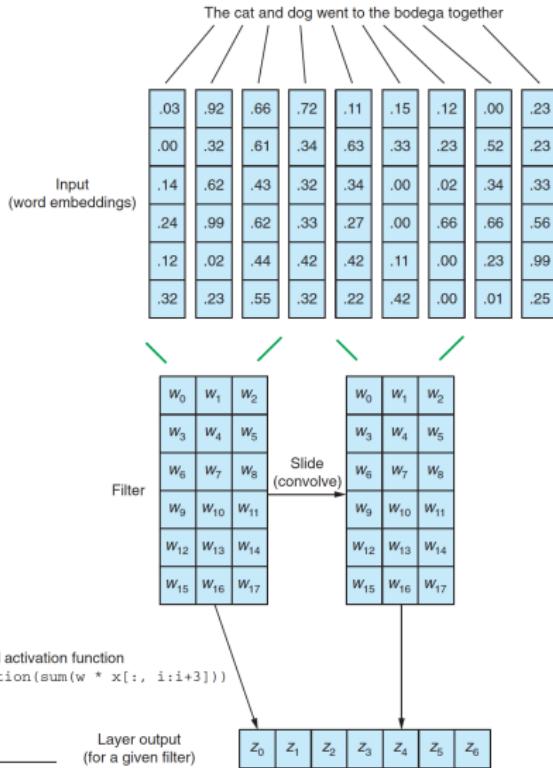
²|l2r or r2l; for some languages it's the vertical direction that matters (e.g., Japanese)
(Lane et al., 2019, p. 229)

But we do have 2D “filters”

Words are represented with word embeddings: vectors

But we do have 2D “filters”

Words are represented with word embeddings: vectors



(Lane et al., 2019, p. 229)

The convolution is (practically) the same as for images

- We now *convolve* in one dimension (not two)

The convolution is (practically) the same as for images

- We now *convolve* in one dimension (not two)
- The computation order is irrelevant, but the outputs have to be fed in the same order

The convolution is (practically) the same as for images

- We now *convolve* in one dimension (not two)
- The computation order is irrelevant, but the outputs have to be fed in the same order
- The filters' weights are fixed for a full sample (parallel computation)

The convolution is (practically) the same as for images

- We now *convolve* in one dimension (not two)
- The computation order is irrelevant, but the outputs have to be fed in the same order
- The filters' weights are fixed for a full sample (parallel computation)
- Their outputs become the features for the classifier

The convolution is (practically) the same as for images

- We now *convolve* in one dimension (not two)
- The computation order is irrelevant, but the outputs have to be fed in the same order
- The filters' weights are fixed for a full sample (parallel computation)
- Their outputs become the features for the classifier

 Let us see

Padding

- (In general) in image processing the inputs are of fixed size, regardless of the instance (same source!)

Padding

- (In general) in image processing the inputs are of fixed size, regardless of the instance (same source!)
- Texts are not fixed length (regardless of their source)

Padding

- (In general) in image processing the inputs are of fixed size, regardless of the instance (same source!)
- Texts are not fixed length (regardless of their source)
- Instances longer than `maxlen` will be truncated

Padding

- (In general) in image processing the inputs are of fixed size, regardless of the instance (same source!)
- Texts are not fixed length (regardless of their source)
- Instances longer than `maxlen` will be truncated
- Instances shorter than `maxlen` will be **padded**

Padding

- (In general) in image processing the inputs are of fixed size, regardless of the instance (same source!)
- Texts are not fixed length (regardless of their source)
- Instances longer than `maxlen` will be truncated
- Instances shorter than `maxlen` will be **padded**

$x_0, x_1, x_2, x_3, \dots x_{398} x_{399} x_{400} x_{401}$

Padding

- (In general) in image processing the inputs are of fixed size, regardless of the instance (same source!)
- Texts are not fixed length (regardless of their source)
- Instances longer than `maxlen` will be truncated
- Instances shorter than `maxlen` will be **padded**

$x_0, x_1, x_2, x_3, \dots x_{398} x_{399} x_{400} x_{401}$

$x_0, x_1, x_2, x_3, \dots x_{397}$ PAD PAD

Padding

- (In general) in image processing the inputs are of fixed size, regardless of the instance (same source!)
- Texts are not fixed length (regardless of their source)
- Instances longer than `maxlen` will be truncated
- Instances shorter than `maxlen` will be **padded**

$x_0, x_1, x_2, x_3, \dots x_{398} x_{399} x_{400} x_{401}$

$x_0, x_1, x_2, x_3, \dots x_{397}$ PAD PAD

 Let us see

Pooling

- For each filter one new version of the instance is produced (250 in the example)

Pooling

- For each filter one new version of the instance is produced (250 in the example)
- Pooling evenly divides the output of each filter into subsections

Pooling

- For each filter one new version of the instance is produced (250 in the example)
- Pooling evenly divides the output of each filter into subsections
- It selects (or computes) a representative value for each subsection

Pooling

Pooling is “the CNN path to dimensionality reduction [...] by learning higher-order representations of the source data” (Lane et al., 2019, p. 236)

(Lane et al., 2019, p. 237)

Pooling

Pooling is “the CNN path to dimensionality reduction [...] by learning higher-order representations of the source data” (Lane et al., 2019, p. 236)

- The filters job is finding patterns → relationships between words and their neighbours

Pooling

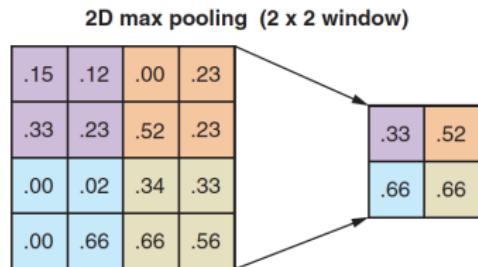
Pooling is “the CNN path to dimensionality reduction [...] by learning higher-order representations of the source data” (Lane et al., 2019, p. 236)

- The filters job is finding patterns → relationships between words and their neighbours
- Pooling in text: a 1D window (e.g., 1×2 or 1×3)

Pooling

Pooling is “the CNN path to dimensionality reduction [...] by learning higher-order representations of the source data” (Lane et al., 2019, p. 236)

- The filters job is finding patterns → relationships between words and their neighbours
- Pooling in text: a 1D window (e.g., 1×2 or 1×3)

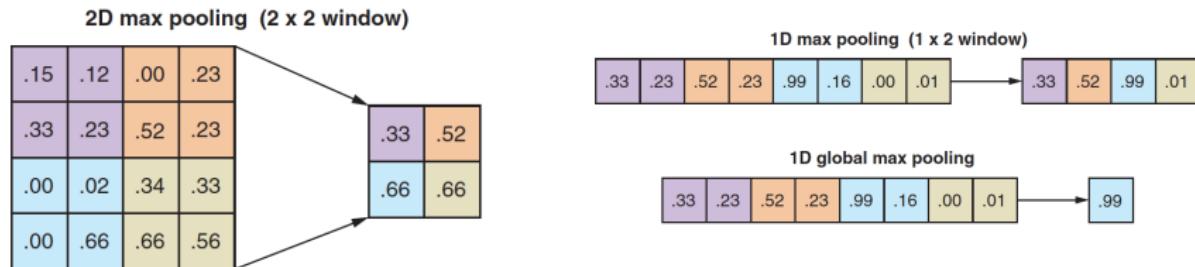


(Lane et al., 2019, p. 237)

Pooling

Pooling is “the CNN path to dimensionality reduction [...] by learning higher-order representations of the source data” (Lane et al., 2019, p. 236)

- The filters job is finding patterns → relationships between words and their neighbours
- Pooling in text: a 1D window (e.g., 1×2 or 1×3)



(Lane et al., 2019, p. 237)

Pooling

Max vs Average Pooling

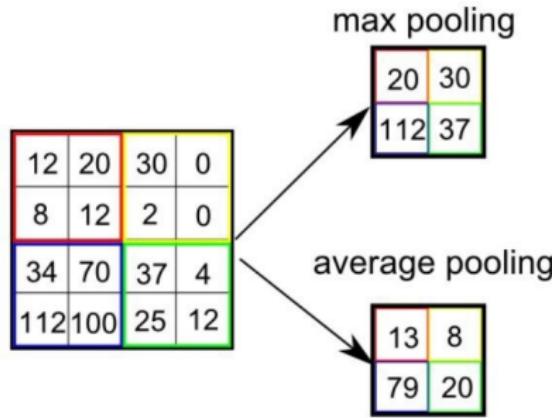


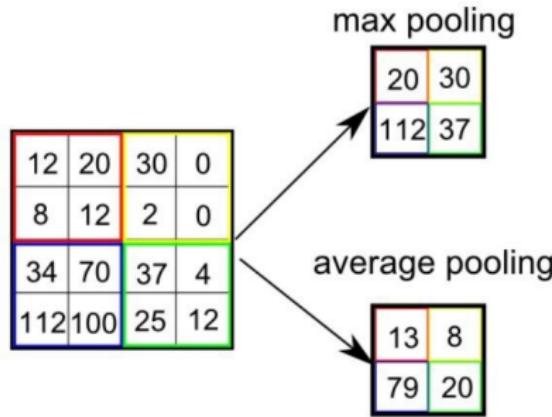
Image borrowed from

www.quora.com/What-is-max-pooling-in-convolutional-neural-networks



Pooling

Max vs Average Pooling



- Average is more intuitive: retaining most of the info

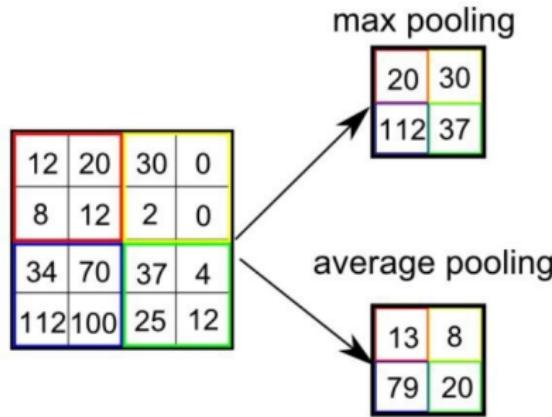
Image borrowed from

www.quora.com/What-is-max-pooling-in-convolutional-neural-networks



Pooling

Max vs Average Pooling



- Average is more intuitive: retaining most of the info
- Max is better: the NN keeps the most prominent feature

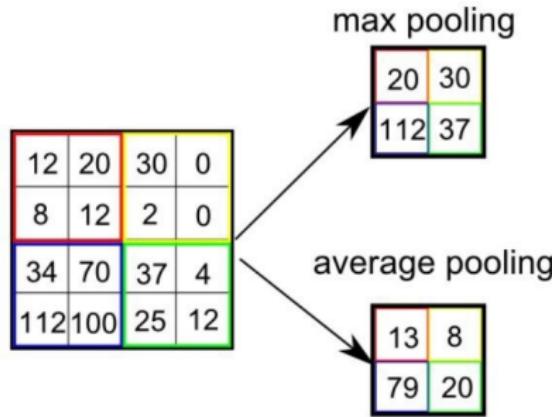
Image borrowed from

www.quora.com/What-is-max-pooling-in-convolutional-neural-networks



Pooling

Max vs Average Pooling



- Average is more intuitive: retaining most of the info
- Max is better: the NN keeps the most prominent feature

Let us see

Image borrowed from

www.quora.com/What-is-max-pooling-in-convolutional-neural-networks



Recap

- Each filter will produce a 1×398 vector

Recap

- Each filter will produce a 1×398 vector
- For each of the 250 filter outputs, we take the single maximum value for each 1D vector

Recap

- Each filter will produce a 1×398 vector
- For each of the 250 filter outputs, we take the single maximum value for each 1D vector
- Output: one 1×250 vector

Recap

- Each filter will produce a 1×398 vector
- For each of the 250 filter outputs, we take the single maximum value for each 1D vector
- Output: one 1×250 vector

This is a crude semantic representation of the text

Dropout: Preventing Overfitting

On each training pass **turn off** a percentage of the input of a layer; it will become 0

- Chosen randomly on each pass

Dropout: Preventing Overfitting

On each training pass **turn off** a percentage of the input of a layer; it will become 0

- Chosen randomly on each pass
- It will not rely heavily on any feature

Dropout: Preventing Overfitting

On each training pass **turn off** a percentage of the input of a layer; it will become 0

- Chosen randomly on each pass
- It will not rely heavily on any feature
- It will generalise better

Dropout: Preventing Overfitting

On each training pass **turn off** a percentage of the input of a layer; it will become 0

- Chosen randomly on each pass
- It will not rely heavily on any feature
- It will generalise better
- Dropout is applied during training only



Photogram from the film "The Platform" (2019)

Dropout: Preventing Overfitting

On each training pass **turn off** a percentage of the input of a layer; it will become 0

- Chosen randomly on each pass
- It will not rely heavily on any feature
- It will generalise better
- Dropout is applied during training only



Let us see

Photogram from the film "The Platform" (2019)

Workhorse Loss Functions

Out of the 13+ available loss functions:

binary_crossentropy: the output neuron is either on or off

Workhorse Loss Functions

Out of the 13+ available loss functions:

`binary_crossentropy`: the output neuron is either on or off

`categorical_crossentropy`: the output is one out of many classes

Workhorse Loss Functions

Out of the 13+ available loss functions:

`binary_crossentropy`: the output neuron is either on or off

`categorical_crossentropy`: the output is one out of many classes

 Let us see

Closing Remarks

- Your input is a series of max 400 words; 300 elements each
- Nothing prevents you from stacking other embeddings (think of RGB)

Image borrowed from <https://blog.mapillary.com>



Closing Remarks

- Your input is a series of max 400 words; 300 elements each
- Nothing prevents you from stacking other embeddings (think of RGB)
- The output of the convolution layer is tied to the task (in this case, sentiment analysis)

Closing Remarks

- Your input is a series of max 400 words; 300 elements each
- Nothing prevents you from stacking other embeddings (think of RGB)
- The output of the convolution layer is tied to the task (in this case, sentiment analysis)
- A CNN is more efficient, thanks to the pooling process and the filters

Closing Remarks

- Your input is a series of max 400 words; 300 elements each
- Nothing prevents you from stacking other embeddings (think of RGB)
- The output of the convolution layer is tied to the task (in this case, sentiment analysis)
- A CNN is more efficient, thanks to the pooling process and the filters
- You can add many convolution layers

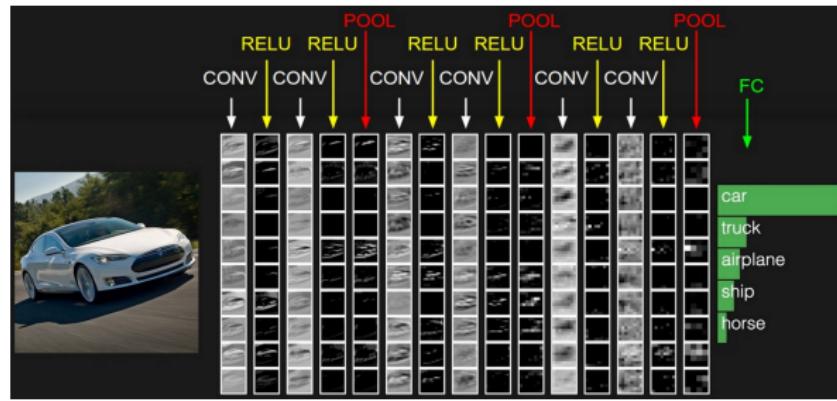


Image borrowed from <https://blog.mapillary.com>

Next time

- Recurrent Neural Networks

References

Lane, H., C. Howard, and H. Hapkem

2019. *Natural Language Processing in Action*. Shelter Island, NY:
Manning Publication Co.