ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA
CAMPUS DI FORLÌ

91258 / B0385
Natural Language Processing

Lesson 10. "One" Neuron

Alberto Barrón-Cedeño
a.barron@unibo.it

29/10/2025

Previously

- From Words to Topics

- ML pipeline

Table of Contents

Chapter 5 of Lane et al. (2019)

There Was Life Before Deep Learning

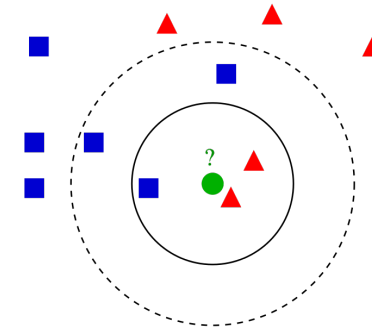Inspired by Hermann Ney's lesson at the 2017 DL Summer School (Bilbao)

There Was Life Before Deep Learning
(And Many Non-NN in-Production Models Prevail)

- Naïve Bayes
- *k*-nearest neighbors
- Random forests
- Support vector machines
- HMM
- Logistic Regression
- . . .

There Was Life Before Deep Learning
k-Nearest Neighbours

The class of ● is the same as the most frequent among its *k* neighbours

https://en.wikipedia.org/wiki/K-nearest_neighbors_algorithm

There Was Life Before Deep Learning
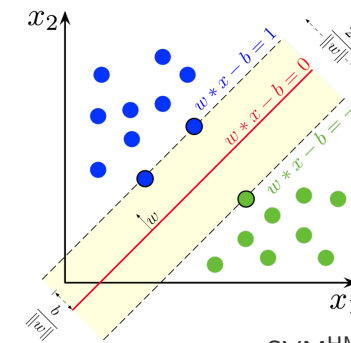Random Forests (showing only one decision tree here)

Playing Golf

Multiple decision trees are learned and the final class is the mode
Picture from https://medium.com/@MrBam44/decision-trees-91f61a42c724
https://en.wikipedia.org/wiki/Random_forest

There Was Life Before Deep Learning
Support Vector Machines

Kernels
- Linear
- Polynomial
- RBF
- Tree

- SVM[HMM] for sequences[a]
- SVM-Rank for ranking
- SVR for regression

[a]Also HMM

There Was Life Before Deep Learning

There are many, many others

- Often they are SotA (or close)
- In general, they are *cheaper*
- In general, they require *less* data
- Some of them are *explainable*
- In general, representations have to be *engineered*

Some History

Some History

Opening paragraph of Rosenblatt (1957)'s The Perceptron—a perceiving and recognizing automaton

```
        Since the advent of electronic computers and modern servo
systems, an increasing amount of attention has been focused on the
feasibility of constructing a device possessing such human-like functions
as perception, recognition, concept formation, and the ability to genera-
lize from experience.  In particular, interest has centered on the idea
of a machine which would be capable of conceptualizing inputs impinging
directly from the physical environment of light, sound, temperature, etc.
-- the "phenomenal world" with which we are all familiar -- rather than
requiring  the intervention of a human agent to digest and code the
necessary information.
```

AI Winters

1974–1980	First major winter
1987–1993	Second major winter
1966	Failure of MT
1970	Abandonment of connectionism (explain mental phenomena using artificial neural networks)
1971–75	DARPA's frustration wrt CMU's speech recognition research
1973	Lighthill report decreases AI research in the UK[1]
1973–74	DARPA's cutbacks to academic AI research
1987	Collapse of the LISP machine market
1988	Cancellation of new spending on AI by the Strategic Computing Initiative
1993	Resistance to expert systems deployment and maintenance
1990s	End of the Fifth Generation computer project's original goals[2]

[1] https://en.wikipedia.org/wiki/Lighthill_report
[2] https://en.wikipedia.org/wiki/Fifth_generation_computer

The Perceptron

The Perceptron

- Intended to be a machine able of recognising images
- Rough idea:

  Input: features of an image (small subsections)

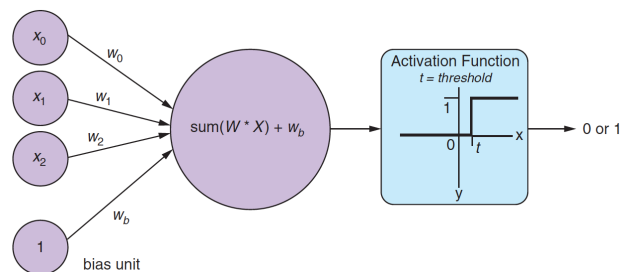 Parameters: weights for each feature (measure of importance)

  Output: *Fire* once all potentiometers pass a certain threshold

  Fired: positive match in the image

  Did not fire: negative class

The Perceptron
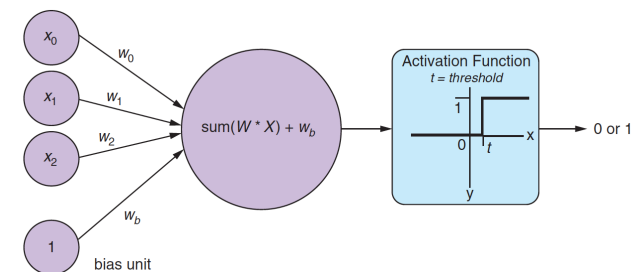Numerical Perceptron[3]

(Lane et al., 2019, p. 158)

- Feature vector: $X = [x_0, x_1, \cdots, x_i, \cdots, x_n]$
- Associated weight (per feature): $W = [w_0, w_1, \cdots, w_i, \ldots, w_n]$
- Sum up: $(x_0 * w_0) + (x_1 * w_1) + \cdots + (x_i * w_i) + \ldots (x_n * w_n)$
- Bias: always-on input (resiliency to inputs of all zeros)
- Activation (step) function

[3]I am intentionally dropping biological references

The Perceptron
Numerical Perceptron

(Lane et al., 2019, p. 158)

$$\hat{y} = f(\vec{x}) = \begin{cases} 1 & \text{if } \sum_{i=0}^{n} x_i w_i > threshold \\ 0 & \text{otherwise} \end{cases} \qquad (1)$$

This perceptron is a special case of a *neuron* —the base unit of a neural network

The Perceptron
Without Bias

"The output [of a perceptron] is a linear function of the input"
(Goodfellow et al., 2016, p. 105)

$$\hat{y} = w^T x \tag{2}$$

The Perceptron
Without Bias

```python
import matplotlib.pyplot as plt
import numpy as np
for i in range(-5, 5, 1):
    fig, ax = plt.subplots(figsize = (5,5))
    ax.spines['left'].set_position('center')
    ax.spines['bottom'].set_position('center')
    ax.spines['right'].set_color('none')
    ax.spines['top'].set_color('none')
    ax.set(title='$y=w^Tx$')
    x = np.arange(-5.0, 5.0, 0.01)
    plt.xlim((-5,+5))
    plt.ylim((-5,+5))
    ax.set(title='$y={}x$'.format(i))
    y = i*x   #1 + np.sin(2 * np.pi * x)
    ax.plot(x, y)
    fig.savefig("linear_w{}.png".format(i))
    plt.show()
```

The Perceptron
Without Bias

Plotting with different values of w; do you see an issue?
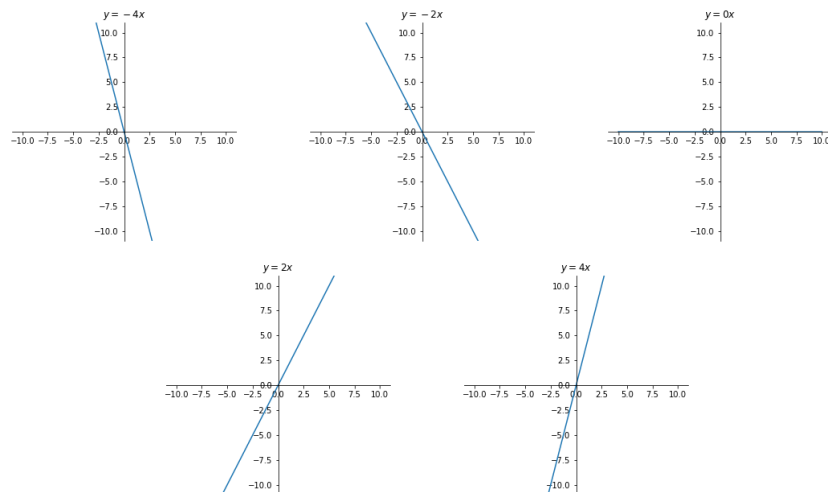
The Perceptron
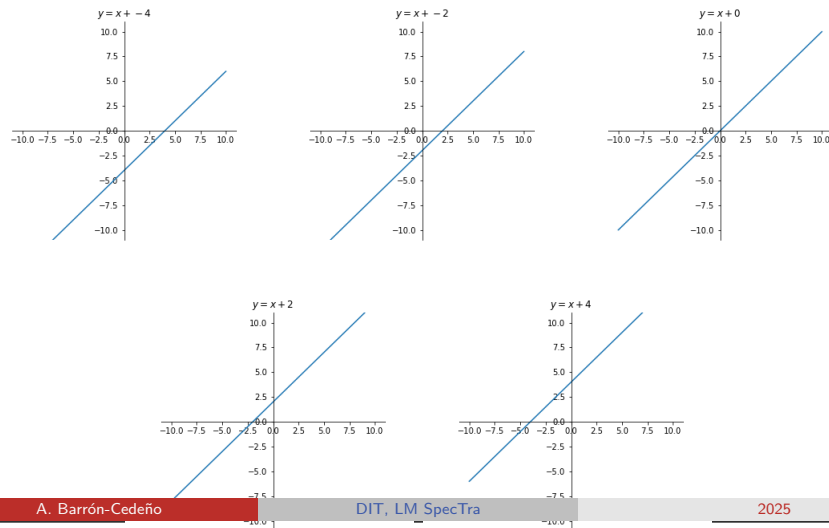### With Bias

$$\hat{y} = w^T x + b \tag{3}$$

"[. . .] the mapping from parameters to predictions is still a linear function but the mapping from features to predictions is now an affine function"
(Goodfellow et al., 2016, p. 107)

(does not need to pass by the origin)

The Perceptron
Without Bias

Plotting with $w = 1$ and different values of b

The Perceptron
Typical Learning Process (1/2)

Given an annotated dataset...

- start with a random weight initialisation from a normal distribution

$$\vec{w} \sim \mathcal{N}(\mu, \sigma^2) \text{ with } \mu \sim 0 \text{ (but do not use 0!)}$$

- feed one instance and see if the predicted class is correct

```
1: if the class is correct then
2:     do nothing
3: else
4:     adjust the weights (slightly; not until getting the class right!)
```

Each weight is adjusted by how much it contributed to the resulting error

---

The Perceptron
Typical Learning Process (2/2)

- All instances in the training data are fed a number of times (iterations): epoch

- Typical stop criteria include
  - $error < \epsilon$ (convergence)
  - $error$ stabilisation
  - max number of epochs reached

The Perceptron
Example 1: Logical OR function

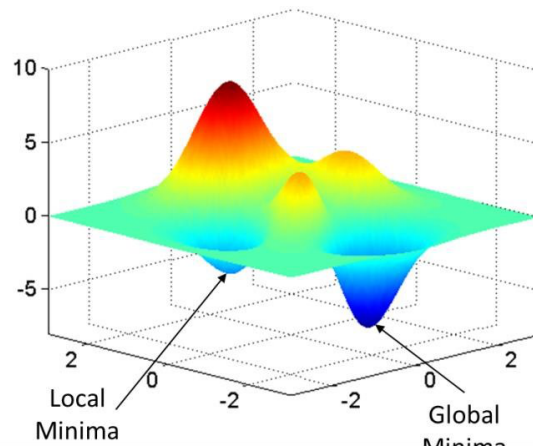| input		output
0	0	0
0	1	1
1	0	1
1	1	1

▤ Let us see

Mr. Perceptron can learn!

This learning model is called linear regression (another ML alternative)

The Perceptron
Drawback: Local vs Global Minimum
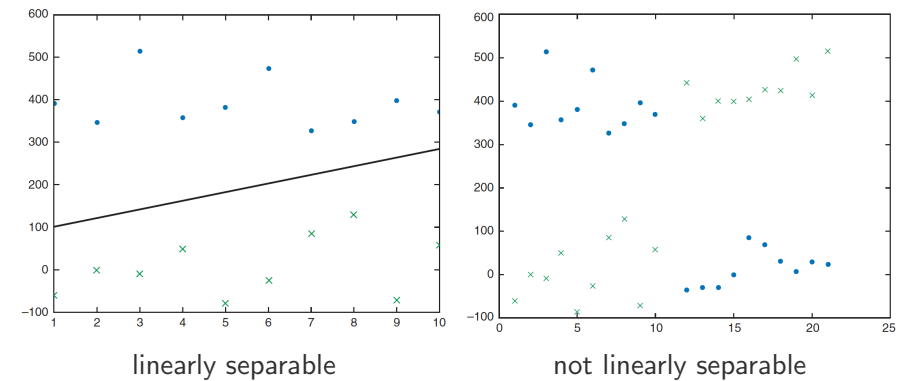

Local Minima

Global Minima

No guarantee that the model will reach the global optimal solution

The Perceptron
Drawback: Linearly separable

The perceptron can only deal with linearly separable data

linearly separable   not linearly separable

Plots from (Lane et al., 2019, p. 164–165)

The Perceptron
Example 2: Logical XOR

We have learned a logical OR function . . .

Can we learn a logical XOR?

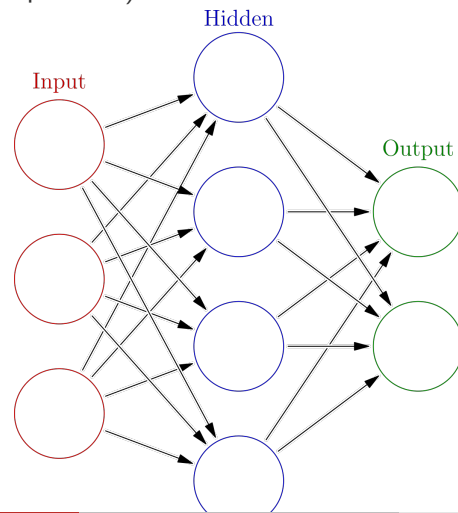| input		output
0	0	0
0	1	1
1	0	1
1	1	0

📖 Let us see

Mr. Perceptron cannot learn!   . . . winter

More than One Neuron

Neural Networks

A neural network is a combination of multiple perceptrons (and it can deal with more complex patterns)

Fully-connected neural network

Some Formalisms

Input $x = [x_1, x_2, x_3, \ldots, x_k]$
Output $f(x)^4$
Answer y

Cost Function[5] Quantifier of the mismatch between actual and predicted output

$$err(x) = |y - f(x)| \qquad (4)$$

Training goal Minimising the cost function across all input samples

$$J(x) = \min \sum_{i=1}^{n} err(x_i) \qquad (5)$$

---

[4]aka $\hat{y}$
[5]aka loss function

Next

- Backpropagation (briefly)
- Activation functions
- Keras

References

Goodfellow, I., Y. Bengio, and A. Courville
2016. *Deep Learning*. MIT Press. http://www.deeplearningbook.org.

Lane, H., C. Howard, and H. Hapkem
2019. *Natural Language Processing in Action*. Shelter Island, NY: Manning Publication Co.

Rosenblatt, F.
1957. The perceptron—a perceiving and recognizing automaton. Technical Report 85-460-1, Cornell Aeronautical Laboratory, Buffallo, NY.