



ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA
CAMPUS DI FORLÌ

91258 / B0385

Natural Language Processing

Lesson 13. Hands on Word Embeddings

Alberto Barrón-Cedeño
a.barron@unibo.it

10/11/2025

Previously

- Skip-gram
- CBOW

Table of Contents

1. Pre-Trained Models
2. Gensim
3. Model Construction
4. GloVe
5. fastText

Chapter 6 of Lane et al. (2019)

Pre-Trained Models

Some Pre-Trained Models

Model	Provider	Description
word2vec	Google	300D from English Google News articles ¹
fastText	Facebook	157 languages from Wikipedia and Common Crawl ²
word2vec/GloVe	CNR	Italian embeddings from the Wikipedia
word2vec	UCampania	Italian embeddings ³

There are many pre-trained models and diverse libraries to handle them.
Just query your favorite search engine

¹<https://drive.google.com/file/d/0B7XkCwpI5KDYN1NUTT1SS21pQmM>


²<https://fasttext.cc>

³<https://mlunicampania.gitlab.io/italian-word2vec/>

Gensim

Gensim

- Scalable, open source, and efficient Python library
- It includes many resources, including word2vec, doc2vec, FastText, LDA, and more
- All information, including (very nice) documentation at <https://radimrehurek.com/gensim/>

 Let us see


Gensim

Most similar items

```
word_vectors.most_similar()
```

Among the most interesting arguments:

- positive** list of vectors to be added together before looking for the neighbours
- negative** subtraction (or exclusion) of the elements
- topn** number of elements to retrieve


 Let us see

Gensim

Least similar items (closed set)

```
word_vectors.doesnt_match()
```

It returns the element from the input list with the lowest similarity with respect to the rest

 Let us see

Gensim

More operations

Adding and Subtracting

We can use `most_similar()` again, this time with the negative parameter

 Let us see

Computing similarities

```
word_vectors.similarity()
```

 Let us see

Gensim

Getting the Vectors

Gensim (and other libraries) have implemented these interfaces to perform some *standard* operations

To go beyond, one needs to get access to the actual vectors

```
word_vectors[word]
```

 Let us see

Model Construction

Model Construction

Considerations

- If you are working in other language than English, Google's provided word2vec is not an option (FastText might be)
- Google's word2vec is built on news; fastText has versions built on the Wikipedia and on common crawl. . . analysing scientific papers or literature?
Probably not
- You want to work on COVID-19 or any other recent topic?
Many relevant terms **wont appear**

Alternatives

- Opting for some of the previous representations
- Getting a more up-to-date existing embedding space
- **Build your own model**

Model Construction

Pre-Processing

Typical pre-processing pipeline

- Tokenisation
- Lowercasing (optional)
- Sentence splitting

Input Embedded list of tokenised sentences


$$[[w_{0,0} \ w_{0,1} \ w_{0,2} \ \dots \ w_{0,k}], [w_{1,0} \ w_{1,1} \ w_{1,2} \ \dots \ w_{1,l}], \dots [w_{x,0} \ w_{x,1} \ \dots \ w_{x,m}]]$$

Model Construction

Training

Training a word2vec model with gensim

Tutorial: <https://rare-technologies.com/word2vec-tutorial/>

 Let us see

Considerations

- Training on relatively large corpora might take some time
(Brown is small and took me a bit less than 1 minute on a 2.5GHz Quad-Core i7, 16GB RAM)
- Large corpora (e.g., the Wikipedia) can require a significant amount of time/memory

Model Construction

Trimming and Saving

Reminder We do not care about the output

```
model.init_sims(replace=True)
```

- Freezes the model
- Stores the hidden-layer weights
- Discards the output-layer weights

not necessary since gensim 4.0

Now we simply have to save the model with `model.save()`

 Let us see

GloVe

GloVe

Global Vectors (Pennington et al., 2014)⁴

- It uses a global word-word co-occurrence matrix
- Learning objective: word vectors such that their dot product equals the logarithm of the words' probability of co-occurrence
- It produces similar matrices to word2vec
- It converges, even with smaller corpora
- It is more accurate with the same amount of data

⁴<https://nlp.stanford.edu/projects/glove/>

GloVe

GloVe vs word2vec

RaRe Technologies comparison⁵

Settings: 600 dims • context window of 10 • 1.9B words of *en* Wikipedia.

Algorithm	acc (word analogy)*	wallclock time	peak RAM (MB)
I/O only	—	3m	25
GloVe, 10 epochs, lr 0.05	67.1	4h12m	9,414
GloVe, 100 epochs, lr 0.05	67.3	18h39m	9,452
word2vec, hierarchical skip-gram, 1 epoch	57.4	3h10m	266
word2vec, negative sampling (10 samples), 1 epoch	68.3	8h38m	628
word2vec, Google 300d	55.3	—	—

* *a is to b as c is to ?*
⁵rare-technologies.com/making-sense-of-Word2vec/#glove_vs_word2vec

fastText

fastText

Predicts the surrounding **character [2, 3]-grams** rather than the surrounding words (Bojanowski et al., 2017)⁶

- Pre-trained models available in 250+ languages
- Built on Wikipedia editions (variable quality)
- Built on common crawl

Models available at <https://github.com/facebookresearch/fastText/blob/master/docs/crawl-vectors.md>

 Let us see

⁶<https://github.com/facebookresearch/fastText>

Some Remarks

LSA is a better (faster) option for long documents e.g., for clustering

Online learning An existing model can be *adapted* (but new words cannot be added)

doc2vec possible representation based on linear combinations of word2vec

References

- Bojanowski, P., E. Grave, A. Joulin, and T. Mikolov
2017. Enriching Word Vectors with Subword Information. *Transactions of the Association for Computational Linguistics*, 5:135–146.
- Lane, H., C. Howard, and H. Hapkem
2019. *Natural Language Processing in Action*. Shelter Island, NY: Manning Publication Co.
- Pennington, J., R. Socher, and C. Manning
2014. GloVe: Global Vectors for Word Representation. In *Empirical Methods in Natural Language Processing (EMNLP)*, Pp. 1532–1543.