



ALMA MATER STUDIORUM
UNIVERSITÀ DI BOLOGNA
CAMPUS DI FORLÌ

91258 / B0385

Natural Language Processing

Lesson 16. Recurrent Neural Networks

Alberto Barrón-Cedeño
a.barron@unibo.it

19/11/2025

Previously

- CNNs for text

A. Barrón-Cedeño

DIT, LM SpecTra

2025 2 / 17

Table of Contents

1. Introduction
2. Keeping the past in mind
3. RNNs in Keras

Chapter 8 of Lane et al. (2019)

Introduction

A. Barrón-Cedeño

DIT, LM SpecTra

2025 3 / 17

A. Barrón-Cedeño

DIT, LM SpecTra

2025 4 / 17

Introduction

CNNs

- Good for analysing *full* texts (~sentences)
- Words tending to appear close to each other are spotted and play a joint role
- Longer relationships —farther than [3, 4] words are ignored

What is missing?

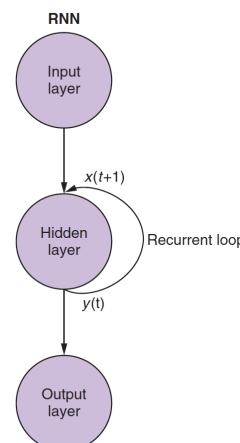
- Keeping track of what happened long ago
- Memory
- Language is **not an image** —no snapshots
- Language is a **sequence**; both text and speech

Keeping the past *in mind*

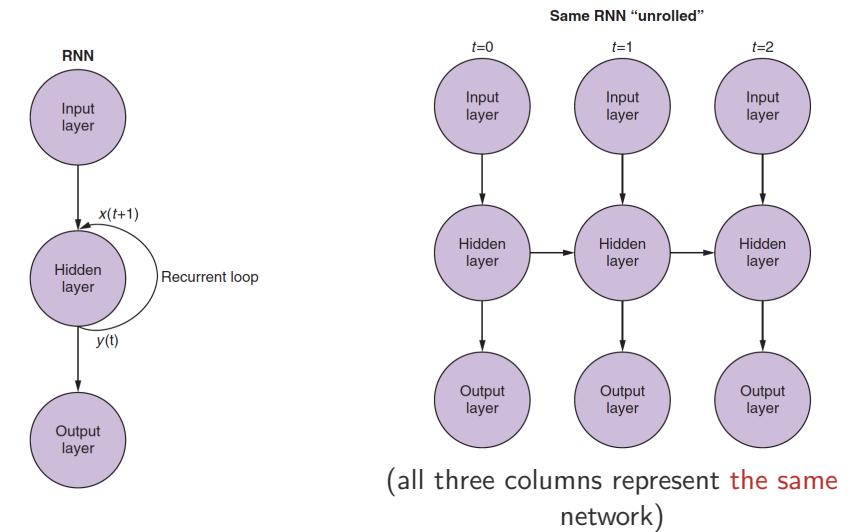
Remembering the Past

$w_0 w_1 w_2 w_3 \dots w_{t-1} w_t w_{t+1}$

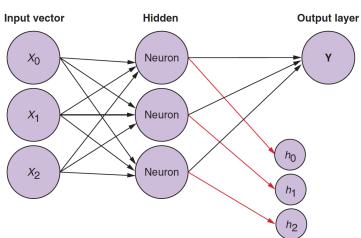
- To understand a text at time t , we need to consider what happened at time $t - k$
- Recurrent neural nets (RNN) come into play
- RNNs combine what happened **before** with what is happening **now**



Full feed-forward networks that consider their own output



Zooming into the unrolled RNN: t and $t + 1$

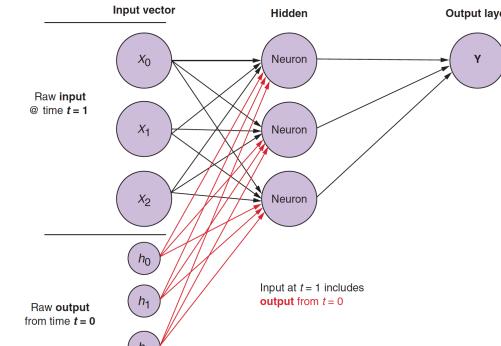


$t = 0$

- The red arrows are just *standard* connections, with weights
- Now we can feed the text, **one word at a time**

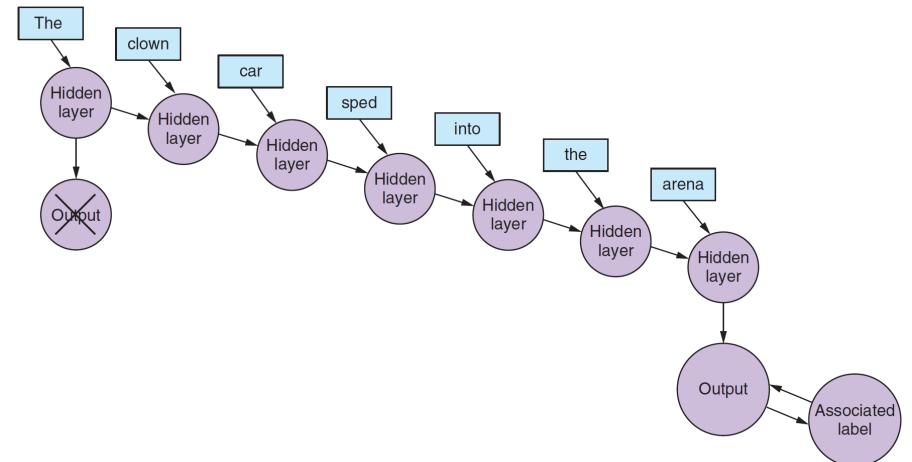
(Lane et al., 2019, p. 252–253)

A. Barrón-Cedeño



$t = 1$

“Multiple inputs, one output”



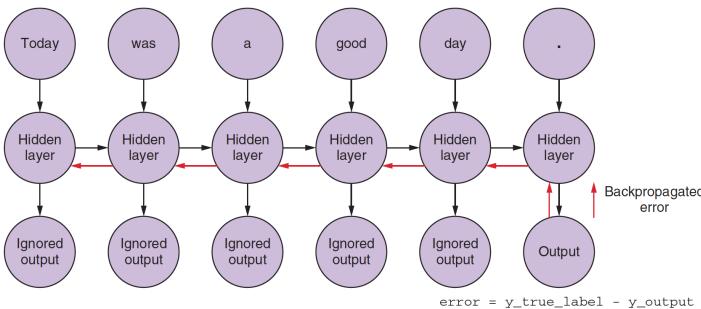
- No more length constraints (although we have to be reasonable)
- No more a bunch of snapshots; there is a sense of time

(Lane et al., 2019, p. 257)

A. Barrón-Cedeño

DIT, LM SpecTra 2025 10 / 17

Backpropagation through Time: the “Vanilla” Way



- All intermediate outputs are ignored; the loss is computed at the end
- The same chain rule is applied to do backpropagation; but this time it heads to “the past”
- The weight corrections are calculated for each t
- The combined updates are applied **only** until reaching $t = 0$

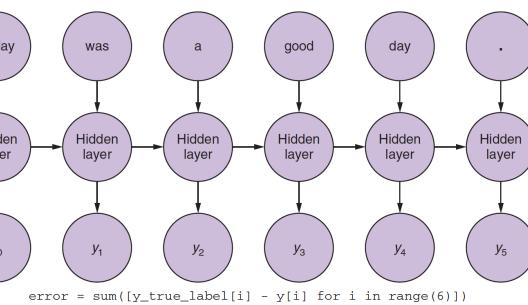
(Lane et al., 2019, p. 256)

A. Barrón-Cedeño

DIT, LM SpecTra

2025 11 / 17

Backpropagation through Time: the Better Way



- We compute the loss combining all intermediate outputs
- The weight corrections are still additive: the update is applied until
 - computing all errors and
 - reaching back to the weight adjustments in $t = 0$



Let us see

(Lane et al., 2019, p. 258)

A. Barrón-Cedeño

DIT, LM SpecTra

2025 12 / 17

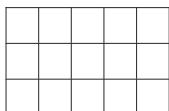
RNNs in Keras

RNN in Keras: further details

- A Dense layer expects a *flat* vector

```
model.add(Flatten())
```

5×3



1×15

+ Flatten →



- In our case: $400 \times 50 \rightarrow 1 \times 20,000$

Let us see

Example derived from

<https://stackoverflow.com/questions/43237124/role-of-flatten-in-keras>

RNN in Keras: what we have so far

We have setup a simple recurrent neural network

- The input sequences have fixed length: 400 tokens (each 300D)
- Our recurrent layer contains 50 units
- The output will be 400×50 :
 - 400 elements
 - one 50D vector each

`return_sequences=True`

`True` return the network value at each t : 400 50D vectors

`False` return a single 50D vector (default)

`True` → this is why we are padding

Let us see

Some parameters are “free”

`embedding_dims` comes from the embedding space; hard to change, but possible: other embeddings, 1-hot

`num_neurons` kind of arbitrary; can be changed

`maxlen` kind of arbitrary; can be changed (or neglected)

`batch_size` bigger→faster (higher local minimum risk)

`epochs` trivial to increase (don't start from scratch each time)

Let us see

Important: unless you have access to HPC, don't go *bananas* when exploring parameters (and perhaps even in that case)

Try some sensitive configurations and keep track of all the settings and outputs¹

¹See, for instance, Fericola et al. (2020) ↗.

References

- Fernicola, F., S. Zhang, F. Garcea, P. Bonora, and A. Barrón-Cedeño
2020. Ariemozione: Identifying emotions in opera verses. In *Italian Conference on Computational Linguistics*.
- Lane, H., C. Howard, and H. Hapkem
2019. *Natural Language Processing in Action*. Shelter Island, NY: Manning Publication Co.